

```
In [1]: ##### Step 1: Import Required Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_m
```

```
In [2]: ##### Step 2: Load and Inspect the Data
df = pd.read_csv("bigml_59c28831336c6604c800002a.csv")
print(df.head())
print(df.info())
print(df.describe())
```

	state	account length	area code	phone number	international plan	\
0	KS	128	415	382-4657	no	
1	OH	107	415	371-7191	no	
2	NJ	137	415	358-1921	no	
3	OH	84	408	375-9999	yes	
4	OK	75	415	330-6626	yes	

	voice mail plan	number vmail messages	total day minutes	total day cal
0	yes	25	265.1	1
10				
1	yes	26	161.6	1
23				
2	no	0	243.4	1
14				
3	no	0	299.4	
71				
4	no	0	166.7	1
13				

```
In [3]: ##### Step 3: Check for Missing Values
print(df.isnull().sum())
```

```
state                0
account length       0
area code            0
phone number         0
international plan   0
voice mail plan      0
number vmail messages 0
total day minutes    0
total day calls      0
total day charge     0
total eve minutes    0
total eve calls      0
total eve charge     0
total night minutes  0
total night calls    0
total night charge   0
total intl minutes   0
total intl calls     0
total intl charge    0
customer service calls 0
churn                0
dtype: int64
```

```
In [4]: ##### Step 4: Data Preprocessing
# Drop unnecessary columns
df.drop(columns=['phone number'], inplace=True)

# Convert target column to numerical
df['churn'] = df['churn'].astype(int)

# Convert categorical columns to numerical
df['international plan'] = df['international plan'].map({'yes': 1, 'no': 0})
df['voice mail plan'] = df['voice mail plan'].map({'yes': 1, 'no': 0})

# One-hot encoding for 'state'
df = pd.get_dummies(df, columns=['state'], drop_first=True)
```

```
In [5]: ##### Step 5: Define Features and Target
y = df['churn']
X = df.drop(columns=['churn'])
```

```
In [6]: ##### Step 6: Split Data into Training and Testing Sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [7]: ##### Step 7: Scale the Features
        scaler = StandardScaler()
        X_train_scaled = scaler.fit_transform(X_train)
        X_test_scaled = scaler.transform(X_test)
```


```
In [8]: ##### Step 8: Train Logistic Regression Model
        log_reg = LogisticRegression()
        log_reg.fit(X_train_scaled, y_train)
        y_pred_log = log_reg.predict(X_test_scaled)
```

```
In [9]: ##### Step 9: Evaluate Logistic Regression Model
print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred_log))
print(classification_report(y_test, y_pred_log))

# Confusion Matrix
plt.figure(figsize=(6, 4))
sns.heatmap(confusion_matrix(y_test, y_pred_log), annot=True, fmt='d', cmap="Blues")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - Logistic Regression")
plt.show()
```

Logistic Regression Accuracy: 0.8575712143928036

	precision	recall	f1-score	support
0	0.87	0.97	0.92	566
1	0.58	0.21	0.31	101
accuracy			0.86	667
macro avg	0.73	0.59	0.61	667
weighted avg	0.83	0.86	0.83	667




```
In [10]: ##### Step 10: Train Decision Tree Model
decision_tree = DecisionTreeClassifier(max_depth=5, random_state=42)
decision_tree.fit(X_train, y_train)
y_pred_tree = decision_tree.predict(X_test)
```

```
In [11]: ##### Sstep 11: Evaluate Decision Tree Model
print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred_tree))
print(classification_report(y_test, y_pred_tree))

# Confusion Matrix for Decision Tree
plt.figure(figsize=(6, 4))
sns.heatmap(confusion_matrix(y_test, y_pred_tree), annot=True, fmt='d', cmap="
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - Decision Tree")
plt.show()
```

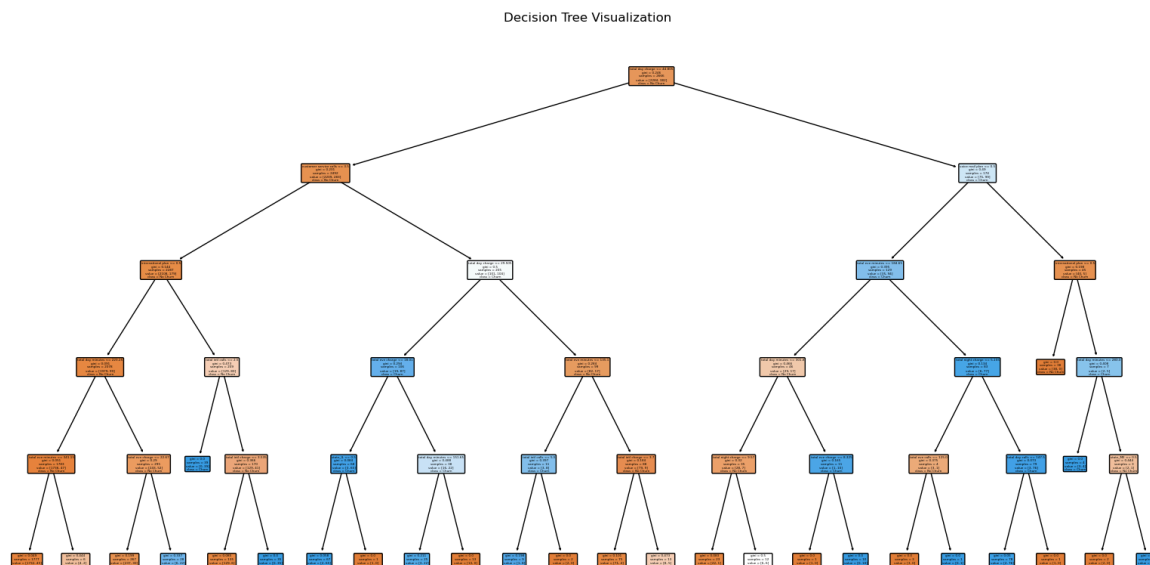
Decision Tree Accuracy: 0.9385307346326837

	precision	recall	f1-score	support
0	0.94	0.99	0.96	566
1	0.91	0.66	0.77	101
accuracy			0.94	667
macro avg	0.92	0.83	0.87	667
weighted avg	0.94	0.94	0.93	667



```
In [13]: from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(20, 10))  
plot_tree(decision_tree, feature_names=X.columns, class_names=['No Churn', 'Churn'])  
plt.title("Decision Tree Visualization")  
plt.show()
```



```
In [14]: # Train a new Decision Tree with a depth Limit
decision_tree = DecisionTreeClassifier(max_depth=4, random_state=42)
decision_tree.fit(X_train, y_train)

# Plot the pruned tree
plt.figure(figsize=(20, 10))
plot_tree(decision_tree, feature_names=X.columns, class_names=['No Churn', 'Churn'],
plt.title("Pruned Decision Tree (max_depth=4)")
plt.show()
```

Pruned Decision Tree (max_depth=4)

